

RAG Pipeline Step 3: Generation

1. Definition

Generation is the "G" in RAG and is the final step in the pipeline. It involves taking the user's original query and the relevant documents (the "context") provided by the **Retrieval** step, and "augmenting" a prompt with them. This augmented prompt is then sent to a Large Language Model (LLM) to synthesize a single, coherent, and factually grounded answer.

2. Core Process

1. **Augment Prompt:** The system takes a prompt template (e.g., "Answer the user's question based *only* on the following context...") and inserts the user's query into the {input} variable and all the retrieved documents into the {context} variable.
2. **Generate Response:** This complete, augmented prompt is sent to an LLM (e.g., GPT-4, Gemini).
3. **Return Answer:** The LLM's job is to reason over the provided context and generate a natural language answer that directly addresses the user's query. This final text is what the user sees.

3. Tools & Frameworks

In LangChain, this step is built by "chaining" components together, most often using the LangChain Expression Language (LCEL).

Tool / Concept	Explanation (in LangChain)	Reason for Use
ChatPromptTemplate	A template for the prompt. It's the "recipe" that instructs the LLM on how to behave (e.g., "Use the following context..."). It contains placeholders for {context} and {input}.	Grounds the model. This is the most critical part for forcing the LLM to base its answer on the retrieved facts, which prevents hallucinations (Source 1.2).
The LLM (ChatOpenAI, etc.)	The core language model (e.g., gpt-4o, gemini-1.5-pro) that performs the reasoning and text generation.	This is the "brain" of the operation, responsible for synthesizing the final answer (Source 1.6).

create_stuff_documents_chain	A helper function that creates a chain specifically designed to "stuff" a list of documents into the {context} variable of a prompt.	Simplifies "stuffing." This is the most common RAG pattern. This function abstracts away the logic of formatting and combining all the retrieved documents (Source 2.1, 2.4).
create_retrieval_chain	A high-level helper that assembles the complete, basic RAG pipeline. It takes a retriever and a combine_docs_chain (like the one above) as input.	Ties it all together. This function creates a runnable chain that automatically handles the two-step process: 1. Call the retriever. 2. Call the generation chain. It returns a dictionary with the answer, context, and input (Source 3.2).
StrOutputParser	An output parser that takes the LLM's complex output (a ChatMessage object) and converts it into a simple string.	Cleans the output. This is typically the last step in the chain, ensuring the final result is a usable text string.
LCEL (`	`)	The LangChain Expression Language , represented by the pipe operator `